

Airbnb Price Prediction and Insights

Part A

```
[106]: #1. Data Exploration and Preprocessing
      ## Load the Data
```

```
[109]: import pandas as pd
```

```
[112]: # Load the dataset
      df = pd.read_csv('Airbnb_data..csv') # Adjust path if needed
```

```
[115]: # Initial Exploration
      df.head()
      df.info()
      df.describe()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 74111 entries, 0 to 74110
```

```
Data columns (total 29 columns):
```

#	Column	Non-Null Count	Dtype
0	id	74111 non-null	int64
1	log_price	74111 non-null	float64
2	property_type	74111 non-null	object
3	room_type	74111 non-null	object
4	amenities	74111 non-null	object
5	accommodates	74111 non-null	int64
6	bathrooms	73911 non-null	float64
7	bed_type	74111 non-null	object
8	cancellation_policy	74111 non-null	object
9	cleaning_fee	74111 non-null	bool
10	city	74111 non-null	object
11	description	74111 non-null	object
12	first_review	58247 non-null	object
13	host_has_profile_pic	73923 non-null	object
14	host_identity_verified	73923 non-null	object
15	host_response_rate	55812 non-null	object
16	host_since	73923 non-null	object
17	instant_bookable	74111 non-null	object
18	last_review	58284 non-null	object
19	latitude	74111 non-null	float64

```

20 longitude          74111 non-null float64
21 name              74111 non-null object
22 neighbourhood     67239 non-null object
23 number_of_reviews 74111 non-null int64
24 review_scores_rating 57389 non-null float64
25 thumbnail_url     65895 non-null object
26 zipcode           73143 non-null object
27 bedrooms          74020 non-null float64
28 beds              73980 non-null float64
dtypes: bool(1), float64(7), int64(3), object(18)
memory usage: 15.9+ MB

```

```

[115]:
count  7.411100e+04  74111.000000  74111.000000  73911.000000  74111.000000 \
mean    1.126662e+07   4.782069    3.155146    1.235263    38.445958
std     6.081735e+06   0.717394    2.153589    0.582044    3.080167
min     3.440000e+02   0.000000    1.000000    0.000000    33.338905
25%     6.261964e+06   4.317488    2.000000    1.000000    34.127908
50%     1.225415e+07   4.709530    2.000000    1.000000    40.662138
75%     1.640226e+07   5.220356    4.000000    1.000000    40.746096
max     2.123090e+07   7.600402   16.000000    8.000000    42.390437

count  74111.000000  74111.000000  57389.000000  74020.000000 \
mean    -92.397525    20.900568    94.067365    1.265793
std     21.705322    37.828641    7.836556    0.852143
min    -122.511500    0.000000    20.000000    0.000000
25%    -118.342374    1.000000    92.000000    1.000000
50%    -76.996965    6.000000    96.000000    1.000000
75%    -73.954660   23.000000   100.000000    1.000000
max    -70.985047   605.000000   100.000000   10.000000

count  73980.000000
mean    1.710868
std     1.254142
min     0.000000
25%     1.000000
50%     1.000000
75%     2.000000
max     18.000000

```

```

[118]: # Data Cleaning
      ## Check for missing values
      df.isnull().sum()

```

```
[118]: id                0
      log_price           0
      property_type      0
      room_type           0
      amenities           0
      accommodates        0
      bathrooms          200
      bed_type            0
      cancellation_policy  0
      cleaning_fee        0
      city                0
      description         0
      first_review        15864
      host_has_profile_pic 188
      host_identity_verified 188
      host_response_rate  18299
      host_since          188
      instant_bookable    0
      last_review         15827
      latitude            0
      longitude           0
      name                0
      neighbourhood      6872
      number_of_reviews    0
      review_scores_rating 16722
      thumbnail_url       8216
      zipcode            968
      bedrooms            91
      beds               131
      dtype: int64
```

```
[121]: # Handle missing values (replace inplace=True with reassignment)

df['bathrooms'] = df['bathrooms'].fillna(df['bathrooms'].median())
print("Missing 'bathrooms' filled with median:", df['bathrooms'].median())

df['bedrooms'] = df['bedrooms'].fillna(df['bedrooms'].median())
print("Missing 'bedrooms' filled with median:", df['bedrooms'].median())

df['beds'] = df['beds'].fillna(df['beds'].median())
print("Missing 'beds' filled with median:", df['beds'].median())

df['review_scores_rating'] = df['review_scores_rating'].
    ↪fillna(df['review_scores_rating'].median())
print("Missing 'review_scores_rating' filled with median:",
    ↪df['review_scores_rating'].median())
```

```
df['host_response_rate'] = df['host_response_rate'].fillna("0%")
print("Missing 'host_response_rate' filled with default value '0%'")
```

Missing 'bathrooms' filled with median: 1.0
 Missing 'bedrooms' filled with median: 1.0
 Missing 'beds' filled with median: 1.0
 Missing 'review_scores_rating' filled with median: 96.0
 Missing 'host_response_rate' filled with default value '0%'

```
[124]: # Fill missing values first (if not already done)
df['host_response_rate'] = df['host_response_rate'].fillna('0%')

# Convert to string, remove '%', convert to float, divide by 100
df['host_response_rate'] = df['host_response_rate'].astype(str).rstrip('%').
    ↪astype(float) / 100.0

print("Converted 'host_response_rate' to numeric format (0.0 to 1.0 range).")
```

Converted 'host_response_rate' to numeric format (0.0 to 1.0 range).

```
[127]: # Extract number of amenities by counting items in the string
df['num_amenities'] = df['amenities'].apply(lambda x: len(x.split(',')) if pd.
    ↪notnull(x) else 0)
print("Created new feature 'num_amenities' by counting amenities per listing.")
```

Created new feature 'num_amenities' by counting amenities per listing.

```
[130]: # Drop only the columns that exist in the DataFrame
columns_to_drop = ['id', 'name', 'description', 'thumbnail_url']
existing_columns_to_drop = [col for col in columns_to_drop if col in df.columns]

df = df.drop(columns=existing_columns_to_drop)
print(f"Dropped columns: {existing_columns_to_drop}")
```

Dropped columns: ['id', 'name', 'description', 'thumbnail_url']

```
[133]: # Define the boolean columns
bool_cols = ['cleaning_fee', 'host_has_profile_pic', 'host_identity_verified',
    ↪'instant_bookable']

# Map 't' to 1 and 'f' to 0
for col in bool_cols:
    df[col] = df[col].map({'t': 1, 'f': 0, True: 1, False: 0})
    print(f"Converted column '{col}' to integers (1 for True/'t', 0 for False/
    ↪'f').")
```

Converted column 'cleaning_fee' to integers (1 for True/'t', 0 for False/'f').
 Converted column 'host_has_profile_pic' to integers (1 for True/'t', 0 for

```
False/'f').
Converted column 'host_identity_verified' to integers (1 for True/'t', 0 for
False/'f').
Converted column 'instant_bookable' to integers (1 for True/'t', 0 for
False/'f').
```

```
[136]: # 2.Model Development
      ## Split the Data
```

```
[139]: from sklearn.model_selection import train_test_split
      from xgboost import XGBRegressor
```

```
[142]: # Step 1: Drop non-numeric or irrelevant object columns
columns_to_drop = ['amenities', 'first_review', 'last_review', 'host_since',
↳ 'zipcode']
X = df.drop(columns=['log_price'] + columns_to_drop)

# Step 2: One-hot encode categorical columns
X = pd.get_dummies(X, drop_first=True)
print("Applied one-hot encoding to categorical features.")

# Step 3: Target variable
y = df['log_price']

# Step 4: Split data
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)
print(f"Split complete: {X_train.shape[0]} training rows, {X_test.shape[0]}
↳ testing rows.")

# Step 5: Train the model
from xgboost import XGBRegressor
model = XGBRegressor()
model.fit(X_train, y_train)
print("Model training complete.")
```

```
Applied one-hot encoding to categorical features.
Split complete: 59288 training rows, 14823 testing rows.
Model training complete.
```

```
[144]: #3.Model Evaluation (10 Marks)
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
import numpy as np
y_pred = model.predict(X_test)
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
mae = mean_absolute_error(y_test, y_pred)
```

```

r2 = r2_score(y_test, y_pred)
print(f'RMSE: {rmse:.2f}')
print(f'MAE: {mae:.2f}')
print(f'R²: {r2:.2f}')

```

RMSE: 0.39

MAE: 0.28

R²: 0.71

```

[147]: # 4. Insights and Visualization
      ## Feature Importance

```

```

[150]: import matplotlib.pyplot as plt
      import seaborn as sns


```

```

[153]: feature_importances = pd.Series(model.feature_importances_, index=X.columns)
      top_features = feature_importances.sort_values(ascending=False)[:10]

      plt.figure(figsize=(10,6))
      sns.barplot(x=top_features.values, y=top_features.index)
      plt.title("Top 10 Important Features for Price Prediction")
      plt.show()

```



```

[ ]:

```