

## Customer Churn Prediction Part B

### 1 Customer Churn Prediction Project Part (B)

#### 1.0.1 1. Load the Data

```
[11]: import pandas as pd
```

```
# Load the dataset
df = pd.read_csv("Customer_data - customer_data.csv")
df.head()
```

```
[11]:  customerID  gender  SeniorCitizen  Partner  Dependents  tenure  PhoneService  \
0  7590-VHVEG  Female                0      Yes           No         1           No
1  5575-GNVDE   Male                0      No            No        34           Yes
2  3668-QPYBK   Male                0      No            No         2           Yes
3  7795-CFOCW   Male                0      No            No        45           No
4  9237-HQITU   Female              0      No            No         2           Yes
```

```
MultipleLines  InternetService  OnlineSecurity  ...  DeviceProtection  \
0  No phone service            DSL              No ...              No
1                No            DSL              Yes ...             Yes
2                No            DSL              Yes ...              No
3  No phone service            DSL              Yes ...             Yes
4                No      Fiber optic            No ...              No
```

```
TechSupport  StreamingTV  StreamingMovies  Contract  PaperlessBilling  \
0          No           No              No  Month-to-month           Yes
1          No           No              No    One year           No
2          No           No              No  Month-to-month           Yes
3          Yes          No              No    One year           No
4          No           No              No  Month-to-month           Yes
```

```
PaymentMethod  MonthlyCharges  TotalCharges  Churn
0  Electronic check           29.85           29.85  No
1    Mailed check           56.95          1889.50  No
2    Mailed check           53.85           108.15  Yes
3  Bank transfer (automatic)    42.30          1840.75  No
4    Electronic check           70.70           151.65  Yes
```

[5 rows x 21 columns]

## 1.0.2 2. Data Cleaning and Preprocessing

```
[12]: # Convert TotalCharges to numeric and drop missing
df['TotalCharges'] = pd.to_numeric(df['TotalCharges'], errors='coerce')
df = df.dropna(subset=['TotalCharges'])

# Replace SeniorCitizen binary to string
df['SeniorCitizen'] = df['SeniorCitizen'].replace({1: 'Yes', 0: 'No'})

# Drop customerID
df.drop('customerID', axis=1, inplace=True)

# Encode target variable
df['Churn'] = df['Churn'].map({'Yes': 1, 'No': 0})

# One-hot encoding
df_encoded = pd.get_dummies(df, drop_first=True)
df_encoded.head()
```

```
[12]:
```

	tenure	MonthlyCharges	TotalCharges	Churn	gender_Male	\
0	1	29.85	29.85	0	False	
1	34	56.95	1889.50	0	True	
2	2	53.85	108.15	1	True	
3	45	42.30	1840.75	0	True	
4	2	70.70	151.65	1	False	

	SeniorCitizen_Yes	Partner_Yes	Dependents_Yes	PhoneService_Yes	\
0	False	True	False	False	
1	False	False	False	True	
2	False	False	False	True	
3	False	False	False	False	
4	False	False	False	True	

	MultipleLines_No phone service	...	StreamingTV_No internet service	\
0	True	...	False	
1	False	...	False	
2	False	...	False	
3	True	...	False	
4	False	...	False	

	StreamingTV_Yes	StreamingMovies_No internet service	StreamingMovies_Yes	\
0	False	False	False	
1	False	False	False	
2	False	False	False	
3	False	False	False	

4	False	False	False
---	-------	-------	-------

	Contract_One year	Contract_Two year	PaperlessBilling_Yes \
0	False	False	True
1	True	False	False
2	False	False	True
3	True	False	False
4	False	False	True

	PaymentMethod_Credit card (automatic)	PaymentMethod_Electronic check \
0	False	True
1	False	False
2	False	False
3	False	False
4	False	True

	PaymentMethod_Mailed check
0	False
1	True
2	True
3	False
4	False

[5 rows x 31 columns]

### 1.0.3 3. Train-Test Split and Feature Scaling

```
[13]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

X = df_encoded.drop('Churn', axis=1)
y = df_encoded['Churn']

# Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

# Scale
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

#### 1.0.4 4. Train Logistic Regression and Random Forest Models

```
[14]: from sklearn.linear_model import LogisticRegression
      from sklearn.ensemble import RandomForestClassifier

      # Use scaled data for logistic regression
      log_model = LogisticRegression(max_iter=2000, solver='lbfgs')
      log_model.fit(X_train_scaled, y_train)

      # Use original data for random forest
      rf_model = RandomForestClassifier(n_estimators=100, random_state=42)
      rf_model.fit(X_train, y_train)
```

```
[14]: RandomForestClassifier(random_state=42)
```

#### 1.0.5 5. Evaluate Model Performance

```
[15]: from sklearn.metrics import classification_report, confusion_matrix
      import matplotlib.pyplot as plt
      import seaborn as sns

      # Predict
      log_pred = log_model.predict(X_test_scaled)
      rf_pred = rf_model.predict(X_test)

      # Reports
      print("Logistic Regression Report:\n", classification_report(y_test, log_pred))
      print("Random Forest Report:\n", classification_report(y_test, rf_pred))

      # Confusion Matrix
      cm = confusion_matrix(y_test, rf_pred)
      plt.figure(figsize=(6, 4))
      sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=['No Churn', 'Churn'],
                  yticklabels=['No Churn', 'Churn'])
      plt.title("Random Forest Confusion Matrix")
      plt.ylabel("Actual")
      plt.xlabel("Predicted")
      plt.tight_layout()
      plt.show()
```


Logistic Regression Report:

	precision	recall	f1-score	support
0	0.83	0.89	0.86	1033
1	0.62	0.52	0.56	374
accuracy			0.79	1407
macro avg	0.73	0.70	0.71	1407

weighted avg	0.78	0.79	0.78	1407
--------------	------	------	------	------

Random Forest Report:

	precision	recall	f1-score	support
0	0.83	0.90	0.86	1033
1	0.63	0.48	0.54	374
accuracy			0.79	1407
macro avg	0.73	0.69	0.70	1407
weighted avg	0.77	0.79	0.78	1407



## 1.1 6. Save Final Models and Processed Data

```
[16]: import joblib

joblib.dump(log_model, "logistic_regression_model_scaled.pkl")
joblib.dump(rf_model, "random_forest_model.pkl")
df_encoded.to_csv("processed_customer_data.csv", index=False)
```

Notebook completed with scaling for Logistic Regression